

Phase 1

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.ensemble import
BaggingClassifier, RandomForestClassifier, AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, classification_report
from sklearn import tree
```

Bagging

```
# Define 9 different colors
colors = ['b', 'g', 'r', 'c', 'm', 'y', 'k', 'purple', 'orange']

# Read data from CSV file
df = pd.read_csv('/content/drive/My Drive/1/Dataset1.csv')

X = df.drop('Label', axis=1)
y = df['Label']

# Scatter plot of Feature 1 vs Feature 2
plt.scatter(df['Feature 1'], df['Feature 2'], c=y,
            cmap=plt.cm.get_cmap('rainbow', len(np.unique(y))))
plt.title('Scatter Plot of Feature1 vs Feature2')
plt.xlabel('Feature1')
plt.ylabel('Feature2')
plt.grid(True)
plt.colorbar(ticks=range(len(np.unique(y))), label='Label')
plt.show()

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
            test_size=0.2, random_state=42)

# Number of estimators to evaluate
num_estimators = [1, 2, 5, 10, 50, 100]

train_accuracies = []
test_accuracies = []

# Loop to test different numbers of estimators
for n in num_estimators:
    # Create Bagging model with different numbers of estimators
```

```

bagging_clf =
BaggingClassifier(base_estimator=DecisionTreeClassifier(max_depth=4,
splitter='random'),
                    n_estimators=n,
                    random_state=42)

# Train the model on training data
bagging_clf.fit(X_train, y_train)

# Predict labels for training and testing data
y_train_pred = bagging_clf.predict(X_train)
y_test_pred = bagging_clf.predict(X_test)

# Compute model accuracy
train_accuracy = accuracy_score(y_train, y_train_pred)
test_accuracy = accuracy_score(y_test, y_test_pred)

# Store accuracies in lists
train_accuracies.append(train_accuracy)
test_accuracies.append(test_accuracy)

# Evaluate the model for training and testing data
train_f1 = f1_score(y_train, y_train_pred, average='micro')
test_f1 = f1_score(y_test, y_test_pred, average='micro')

train_precision = precision_score(y_train, y_train_pred,
average='micro')
test_precision = precision_score(y_test, y_test_pred,
average='micro')

train_recall = recall_score(y_train, y_train_pred,
average='micro')
test_recall = recall_score(y_test, y_test_pred, average='micro')

print(f'n_estimators Bagging = {n}:')
print(f'Train Accuracy: {train_accuracy:.2f}, Test Accuracy:
{test_accuracy:.2f}')
print(f'Train Precision: {train_precision:.2f}, Test Precision:
{test_precision:.2f}')
print(f'Train Recall: {train_recall:.2f}, Test Recall:
{test_recall:.2f}')
print(f'Train F1-score: {train_f1:.2f}, Test F1-score:
{test_f1:.2f}')
print()

# Plot Bagging accuracy versus number of estimators
plt.figure(figsize=(10, 6))
plt.plot(num_estimators, train_accuracies, marker='o', label='Train
Accuracy')
plt.plot(num_estimators, test_accuracies, marker='o', label='Test
Accuracy')

```

```

plt.title('Bagging Accuracy vs. Number of Estimators')
plt.xlabel('Number of Estimators')
plt.ylabel('Accuracy')
plt.xticks(num_estimators)
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# Display five random trees from the model
plt.figure(figsize=(20, 10))
for i in range(5):
    plt.subplot(2, 3, i + 1)

    # Plot the tree
    plot_tree(bagging_clf.estimators_[i], filled=True) # Use
    plot_tree to plot the tree
    plt.title(f'Tree Bagging {i + 1}')

plt.tight_layout()
plt.show()

# Plot decision boundaries
plt.figure(figsize=(15, 10))

# Create a mesh grid
x_min, x_max = X['Feature 1'].min() - 1, X['Feature 1'].max() + 1
y_min, y_max = X['Feature 2'].min() - 1, X['Feature 2'].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))

for i in range(5):
    plt.subplot(2, 3, i + 1)

    # Predict the class for each point in the mesh grid
    Z = bagging_clf.estimators_[i].predict(np.c_[xx.ravel(),
    yy.ravel()])
    Z = Z.reshape(xx.shape)

    # Plot decision boundaries
    plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
    len(np.unique(y))))

    # Plot original data points with predicted labels
    y_pred_tree = bagging_clf.estimators_[i].predict(X)
    for label in np.unique(y):
        plt.scatter(X[y_pred_tree == label]['Feature 1'],
        X[y_pred_tree == label]['Feature 2'],
        color=colors[label], label=f'Label {label}')

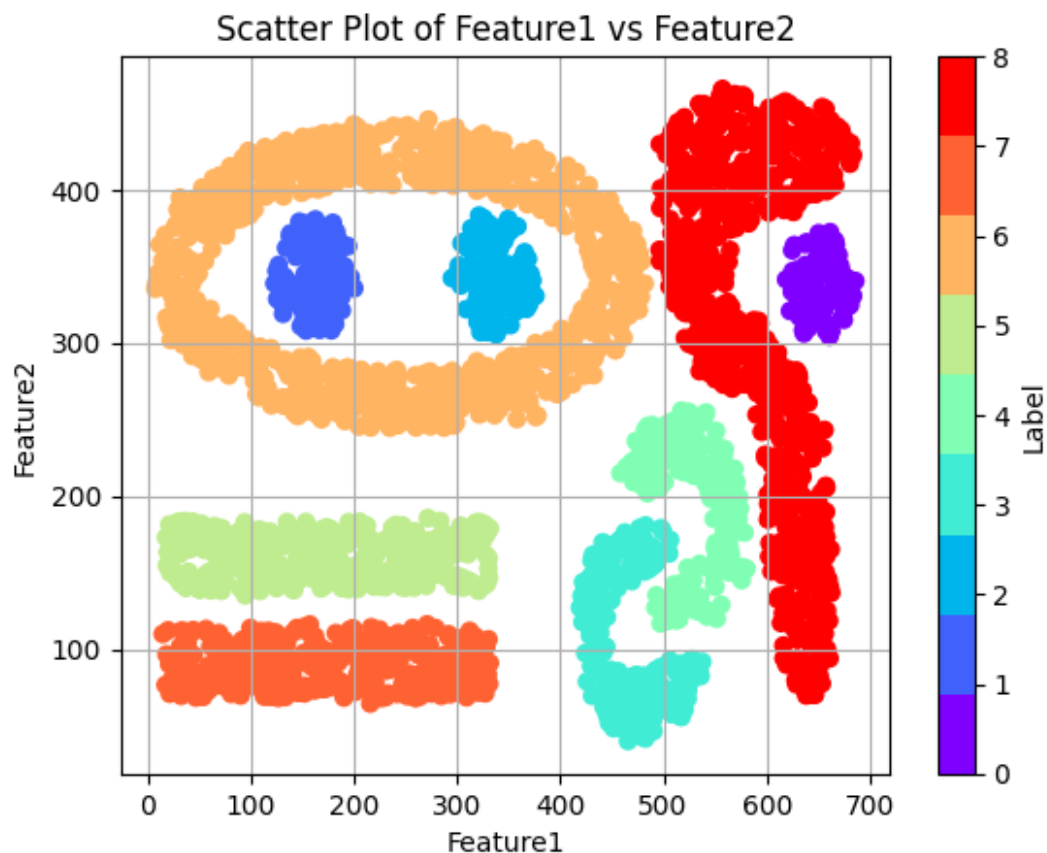
```

```
plt.title(f'Tree Bagging {i + 1}')
plt.legend()
```

```
plt.tight_layout()
plt.show()
```

<ipython-input-54-1fa8d6f22dc9>:11: MatplotlibDeprecationWarning: The get_cmap function was deprecated in Matplotlib 3.7 and will be removed two minor releases later. Use ``matplotlib.colormaps[name]`` or ``matplotlib.colormaps.get_cmap(obj)`` instead.

```
plt.scatter(df['Feature 1'], df['Feature 2'], c=y,
cmap=plt.cm.get_cmap('rainbow', len(np.unique(y))))
```



/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166: FutureWarning: `base_estimator` was renamed to `estimator` in version 1.2 and will be removed in 1.4.

```
warnings.warn(
```

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166: FutureWarning: `base_estimator` was renamed to `estimator` in version 1.2 and will be removed in 1.4.

```
warnings.warn(
```

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:

```
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
warnings.warn(

n_estimators Bagging = 1:
Train Accuracy: 0.81, Test Accuracy: 0.80
Train Precision: 0.81, Test Precision: 0.80
Train Recall: 0.81, Test Recall: 0.80
Train F1-score: 0.81, Test F1-score: 0.80

n_estimators Bagging = 2:
Train Accuracy: 0.83, Test Accuracy: 0.81
Train Precision: 0.83, Test Precision: 0.81
Train Recall: 0.83, Test Recall: 0.81
Train F1-score: 0.83, Test F1-score: 0.81

n_estimators Bagging = 5:
Train Accuracy: 0.85, Test Accuracy: 0.83
Train Precision: 0.85, Test Precision: 0.83
Train Recall: 0.85, Test Recall: 0.83
Train F1-score: 0.85, Test F1-score: 0.83

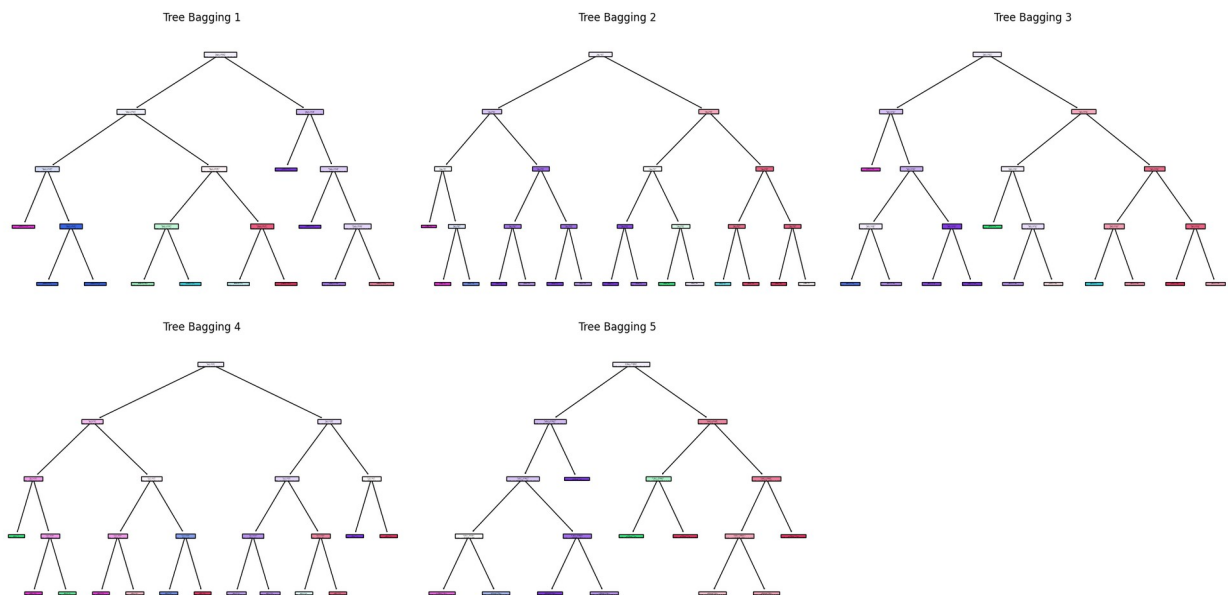

n_estimators Bagging = 10:
Train Accuracy: 0.86, Test Accuracy: 0.85
Train Precision: 0.86, Test Precision: 0.85
Train Recall: 0.86, Test Recall: 0.85
Train F1-score: 0.86, Test F1-score: 0.85

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
warnings.warn(

n_estimators Bagging = 50:
Train Accuracy: 0.87, Test Accuracy: 0.86
Train Precision: 0.87, Test Precision: 0.86
Train Recall: 0.87, Test Recall: 0.86
Train F1-score: 0.87, Test F1-score: 0.86

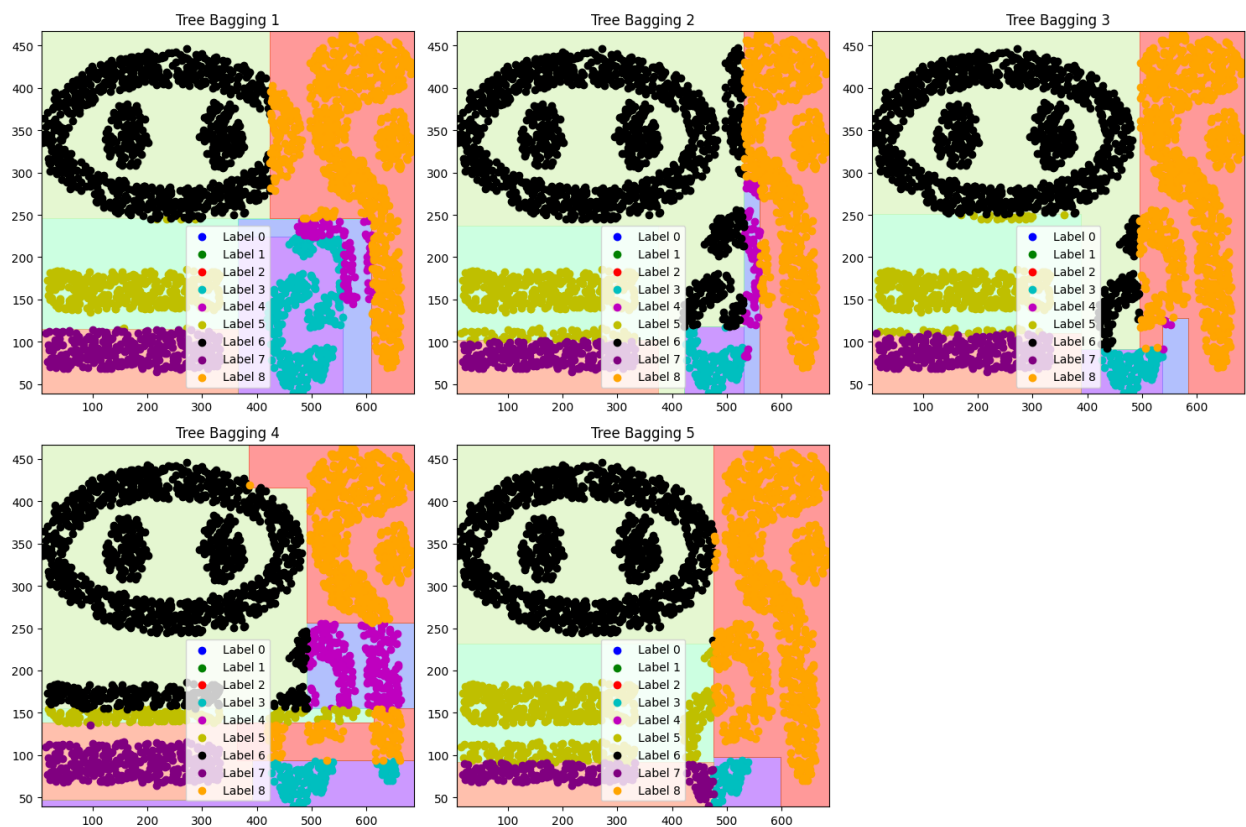
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
warnings.warn(
```

n_estimators Bagging = 100:
Train Accuracy: 0.88, Test Accuracy: 0.87
Train Precision: 0.88, Test Precision: 0.87
Train Recall: 0.88, Test Recall: 0.87
Train F1-score: 0.88, Test F1-score: 0.87



```
<ipython-input-54-1fa8d6f22dc9>:108: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432:
UserWarning: X has feature names, but DecisionTreeClassifier was
fitted without feature names
warnings.warn(
<ipython-input-54-1fa8d6f22dc9>:108: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432:
UserWarning: X has feature names, but DecisionTreeClassifier was
fitted without feature names
warnings.warn(
<ipython-input-54-1fa8d6f22dc9>:108: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432:
UserWarning: X has feature names, but DecisionTreeClassifier was
fitted without feature names
warnings.warn(
<ipython-input-54-1fa8d6f22dc9>:108: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432:
UserWarning: X has feature names, but DecisionTreeClassifier was
```

```
fitted without feature names
warnings.warn(
```



RandomForest

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.tree import plot_tree
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

# Define 9 different colors
colors = ['b', 'g', 'r', 'c', 'm', 'y', 'k', 'purple', 'orange']

# Read data from CSV file
df = pd.read_csv('/content/drive/My Drive/1/Dataset1.csv')

X = df.drop('Label', axis=1)
y = df['Label']

# Scatter plot of Feature 1 vs Feature 2
```



```

plt.scatter(df['Feature 1'], df['Feature 2'], c=y,
            cmap=plt.cm.get_cmap('rainbow', len(np.unique(y))))
plt.title('Scatter Plot of Feature1 vs Feature2')
plt.xlabel('Feature1')
plt.ylabel('Feature2')
plt.grid(True)
plt.colorbar(ticks=range(len(np.unique(y))), label='Label')
plt.show()

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Number of estimators to evaluate
num_estimators = [1, 2, 5, 10, 50, 100]

train_accuracies = []
test_accuracies = []

# Loop to test different numbers of estimators
for n in num_estimators:
    # Create RandomForest model with different numbers of estimators
    rf_clf = RandomForestClassifier(n_estimators=n, max_depth=4,
                                    random_state=42)

    # Train the model on training data
    rf_clf.fit(X_train, y_train)

    # Predict labels for training and testing data
    y_train_pred = rf_clf.predict(X_train)
    y_test_pred = rf_clf.predict(X_test)

    # Compute model accuracy
    train_accuracy = accuracy_score(y_train, y_train_pred)
    test_accuracy = accuracy_score(y_test, y_test_pred)

    # Store accuracies in lists
    train_accuracies.append(train_accuracy)
    test_accuracies.append(test_accuracy)

    # Evaluate the model for training and testing data
    train_f1 = f1_score(y_train, y_train_pred, average='micro')
    test_f1 = f1_score(y_test, y_test_pred, average='micro')

    train_precision = precision_score(y_train, y_train_pred,
                                      average='micro')
    test_precision = precision_score(y_test, y_test_pred,
                                    average='micro')

    train_recall = recall_score(y_train, y_train_pred,

```

```

average='micro')
test_recall = recall_score(y_test, y_test_pred, average='micro')

print(f'n_estimators RandomForest = {n}:')
print(f'Train Accuracy: {train_accuracy:.2f}, Test Accuracy:
{test_accuracy:.2f}')
print(f'Train Precision: {train_precision:.2f}, Test Precision:
{test_precision:.2f}')
print(f'Train Recall: {train_recall:.2f}, Test Recall:
{test_recall:.2f}')
print(f'Train F1-score: {train_f1:.2f}, Test F1-score:
{test_f1:.2f}')
print()

# Plot RandomForest accuracy versus number of estimators
plt.figure(figsize=(10, 6))
plt.plot(num_estimators, train accuracies, marker='o', label='Train
Accuracy')
plt.plot(num_estimators, test accuracies, marker='o', label='Test
Accuracy')
plt.title('RandomForest Accuracy vs. Number of Estimators')
plt.xlabel('Number of Estimators')
plt.ylabel('Accuracy')
plt.xticks(num_estimators)
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# Display five random trees from the model
plt.figure(figsize=(20, 10))
for i in range(5):
    plt.subplot(2, 3, i + 1)

    # Plot the tree
    plot_tree(rf_clf.estimators_[i], filled=True) # Use plot_tree to
plot the tree
    plt.title(f'Tree RandomForest {i + 1}')

plt.tight_layout()
plt.show()

# Plot decision boundaries
plt.figure(figsize=(15, 10))

# Create a mesh grid
x_min, x_max = X['Feature 1'].min() - 1, X['Feature 1'].max() + 1
y_min, y_max = X['Feature 2'].min() - 1, X['Feature 2'].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))

```

```

for i in range(5):
    plt.subplot(2, 3, i + 1)

    # Predict the class for each point in the mesh grid
    Z = rf_clf.estimators_[i].predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    # Plot decision boundaries
    plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))

    # Plot original data points with predicted labels
    y_pred_tree = rf_clf.estimators_[i].predict(X)
    for label in np.unique(y):
        plt.scatter(X[y_pred_tree == label]['Feature 1'],
X[y_pred_tree == label]['Feature 2'],
                    color=colors[label], label=f'Label {label}')

    plt.title(f'Tree RandomForest {i + 1}')
    plt.legend()

plt.tight_layout()
plt.show()


```

<ipython-input-55-27a5bddb699e>:19: MatplotlibDeprecationWarning: The get_cmap function was deprecated in Matplotlib 3.7 and will be removed two minor releases later. Use ``matplotlib.colormaps[name]`` or ``matplotlib.colormaps.get_cmap(obj)`` instead.

```

plt.scatter(df['Feature 1'], df['Feature 2'], c=y,
cmap=plt.cm.get_cmap('rainbow', len(np.unique(y))))

```



```
n_estimators RandomForest = 1:  
Train Accuracy: 0.88, Test Accuracy: 0.85  
Train Precision: 0.88, Test Precision: 0.85  
Train Recall: 0.88, Test Recall: 0.85  
Train F1-score: 0.88, Test F1-score: 0.85
```

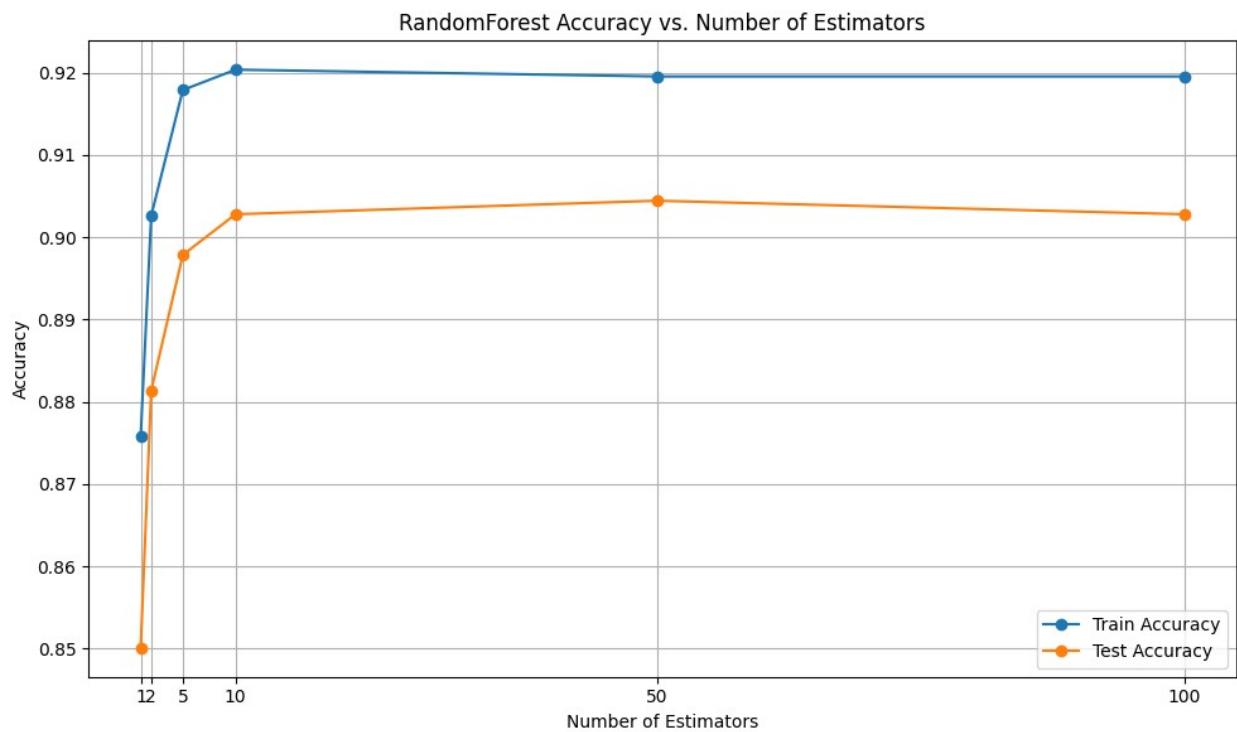
```
n_estimators RandomForest = 2:  
Train Accuracy: 0.90, Test Accuracy: 0.88  
Train Precision: 0.90, Test Precision: 0.88  
Train Recall: 0.90, Test Recall: 0.88  
Train F1-score: 0.90, Test F1-score: 0.88
```

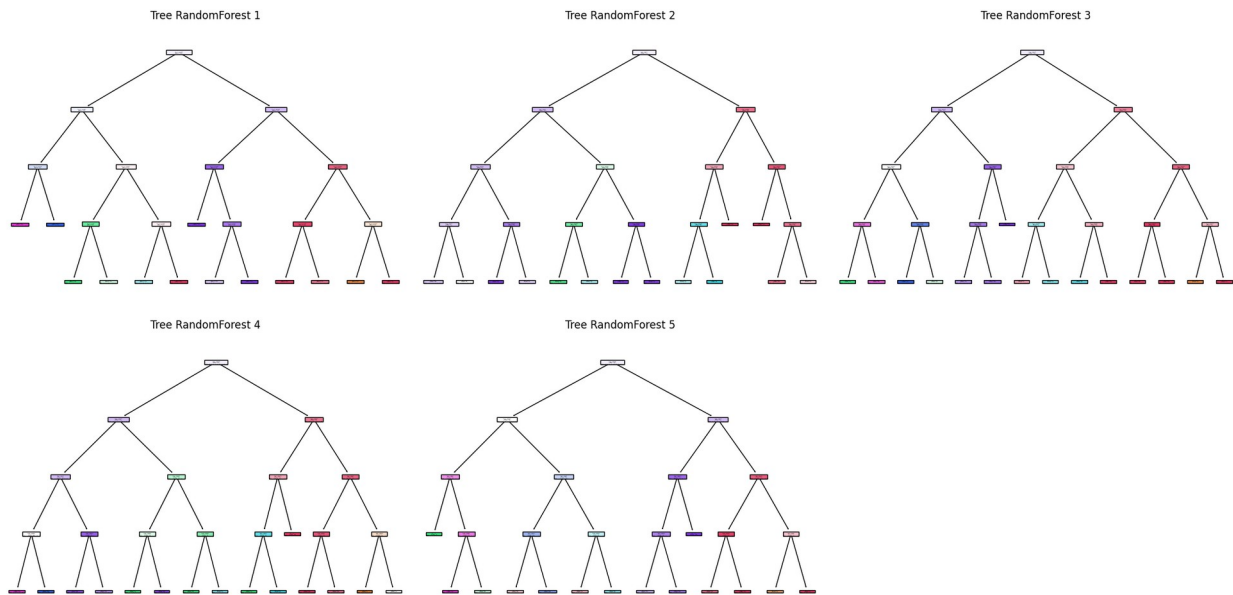
```
n_estimators RandomForest = 5:  
Train Accuracy: 0.92, Test Accuracy: 0.90  
Train Precision: 0.92, Test Precision: 0.90  
Train Recall: 0.92, Test Recall: 0.90  
Train F1-score: 0.92, Test F1-score: 0.90
```

```
n_estimators RandomForest = 10:  
Train Accuracy: 0.92, Test Accuracy: 0.90  
Train Precision: 0.92, Test Precision: 0.90  
Train Recall: 0.92, Test Recall: 0.90  
Train F1-score: 0.92, Test F1-score: 0.90
```

```
n_estimators RandomForest = 50:  
Train Accuracy: 0.92, Test Accuracy: 0.90  
Train Precision: 0.92, Test Precision: 0.90  
Train Recall: 0.92, Test Recall: 0.90  
Train F1-score: 0.92, Test F1-score: 0.90
```

```
n_estimators RandomForest = 100:  
Train Accuracy: 0.92, Test Accuracy: 0.90  
Train Precision: 0.92, Test Precision: 0.90  
Train Recall: 0.92, Test Recall: 0.90  
Train F1-score: 0.92, Test F1-score: 0.90
```



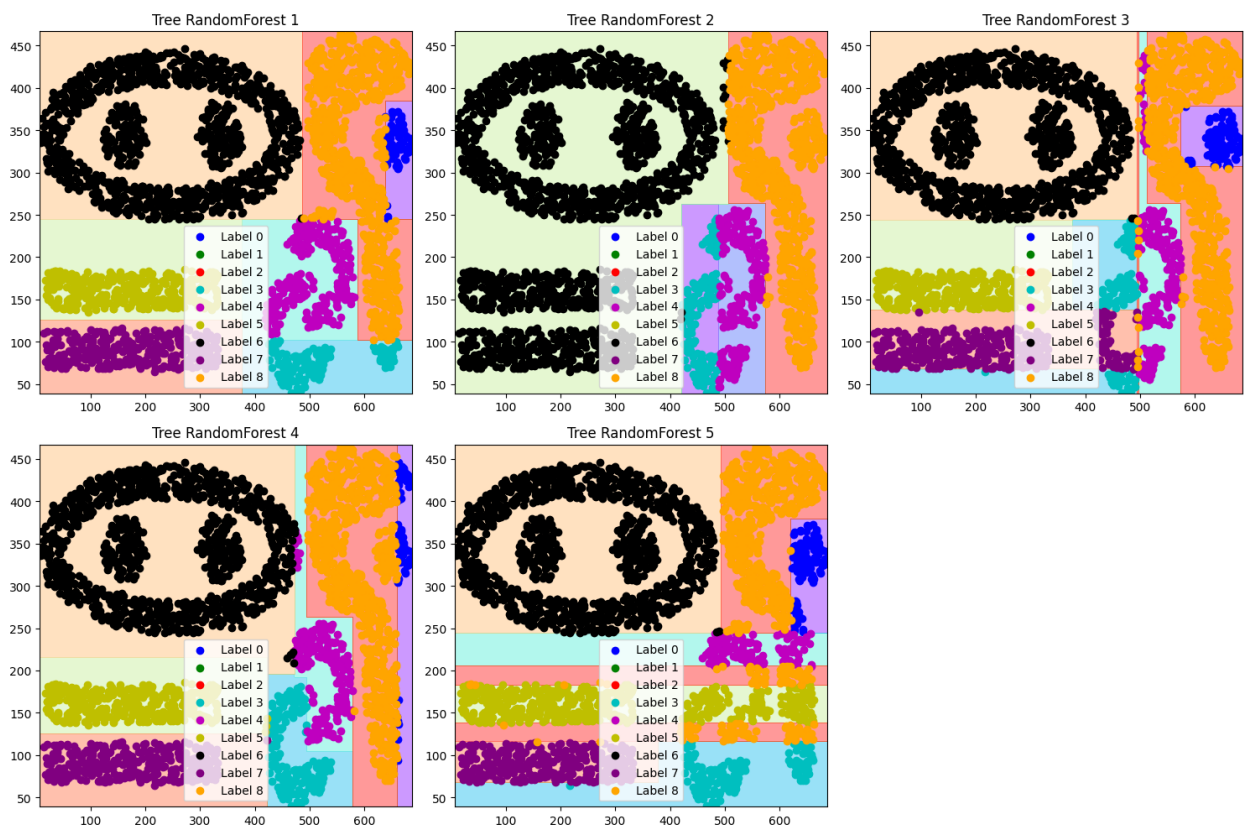


```
<ipython-input-55-27a5bddb699e>:115: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432:
UserWarning: X has feature names, but DecisionTreeClassifier was
fitted without feature names
warnings.warn(
<ipython-input-55-27a5bddb699e>:115: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432:
UserWarning: X has feature names, but DecisionTreeClassifier was
fitted without feature names
warnings.warn(
<ipython-input-55-27a5bddb699e>:115: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432:
UserWarning: X has feature names, but DecisionTreeClassifier was
fitted without feature names
warnings.warn(
```

```

<ipython-input-55-27a5bddb699e>:115: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432:
UserWarning: X has feature names, but DecisionTreeClassifier was
fitted without feature names
warnings.warn(
<ipython-input-55-27a5bddb699e>:115: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432:
UserWarning: X has feature names, but DecisionTreeClassifier was
fitted without feature names
warnings.warn(

```



Adaboost

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import AdaBoostClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

# Define 9 different colors
colors = ['b', 'g', 'r', 'c', 'm', 'y', 'k', 'purple', 'orange']

# Read data from CSV file
df = pd.read_csv('/content/drive/My Drive/1/Dataset1.csv')

X = df.drop('Label', axis=1)
y = df['Label']

# Scatter plot of Feature 1 vs Feature 2
plt.scatter(df['Feature 1'], df['Feature 2'], c=y,
            cmap=plt.cm.get_cmap('rainbow', len(np.unique(y))))
plt.title('Scatter Plot of Feature1 vs Feature2')
plt.xlabel('Feature1')
plt.ylabel('Feature2')
plt.grid(True)
plt.colorbar(ticks=range(len(np.unique(y))), label='Label')
plt.show()

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
            test_size=0.2, random_state=42)

# Number of estimators to evaluate
num_estimators = [1, 2, 5, 10, 50, 100]

train_accuracies = []
test_accuracies = []

# Loop to test different numbers of estimators
for n in num_estimators:
    # Create AdaBoost model with different numbers of estimators
    ada_clf =
AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=4,
            splitter='random'),
                    n_estimators=n, random_state=42)

    # Train the model on training data
    ada_clf.fit(X_train, y_train)
```



```

# Predict labels for training and testing data
y_train_pred = ada_clf.predict(X_train)
y_test_pred = ada_clf.predict(X_test)

# Compute model accuracy
train_accuracy = accuracy_score(y_train, y_train_pred)
test_accuracy = accuracy_score(y_test, y_test_pred)

# Store accuracies in lists
train_accuracies.append(train_accuracy)
test_accuracies.append(test_accuracy)

# Evaluate the model for training and testing data
train_f1 = f1_score(y_train, y_train_pred, average='micro')
test_f1 = f1_score(y_test, y_test_pred, average='micro')

train_precision = precision_score(y_train, y_train_pred,
average='micro')
test_precision = precision_score(y_test, y_test_pred,
average='micro')

train_recall = recall_score(y_train, y_train_pred,
average='micro')
test_recall = recall_score(y_test, y_test_pred, average='micro')

print(f'n_estimators AdaBoost = {n}:')
print(f'Train Accuracy: {train_accuracy:.2f}, Test Accuracy:
{test_accuracy:.2f}')
print(f'Train Precision: {train_precision:.2f}, Test Precision:
{test_precision:.2f}')
print(f'Train Recall: {train_recall:.2f}, Test Recall:
{test_recall:.2f}')
print(f'Train F1-score: {train_f1:.2f}, Test F1-score:
{test_f1:.2f}')
print()

# Plot AdaBoost accuracy versus number of estimators
plt.figure(figsize=(10, 6))
plt.plot(num_estimators, train_accuracies, marker='o', label='Train
Accuracy')
plt.plot(num_estimators, test_accuracies, marker='o', label='Test
Accuracy')
plt.title('AdaBoost Accuracy vs. Number of Estimators')
plt.xlabel('Number of Estimators')
plt.ylabel('Accuracy')
plt.xticks(num_estimators)
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

```

```

# Display five random trees from the model
plt.figure(figsize=(20, 10))
for i in range(5):
    plt.subplot(2, 3, i + 1)

    # Plot the tree
    plot_tree(ada_clf.estimators_[i], filled=True) # Use plot_tree to
    plot the tree
    plt.title(f'Tree AdaBoost {i + 1}')

plt.tight_layout()
plt.show()

# Plot decision boundaries
plt.figure(figsize=(15, 10))

# Create a mesh grid
x_min, x_max = X['Feature 1'].min() - 1, X['Feature 1'].max() + 1
y_min, y_max = X['Feature 2'].min() - 1, X['Feature 2'].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))

for i in range(5):
    plt.subplot(2, 3, i + 1)

    # Predict the class for each point in the mesh grid
    Z = ada_clf.estimators_[i].predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    # Plot decision boundaries
    plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))

    # Plot original data points with predicted labels
    y_pred_tree = ada_clf.estimators_[i].predict(X)
    for label in np.unique(y):
        plt.scatter(X[y_pred_tree == label]['Feature 1'],
X[y_pred_tree == label]['Feature 2'],
                    color=colors[label], label=f'Label {label}')


    plt.title(f'Tree AdaBoost {i + 1}')
    plt.legend()

plt.tight_layout()
plt.show()

```

<ipython-input-56-9d42626d182c>:19: MatplotlibDeprecationWarning: The get_cmap function was deprecated in Matplotlib 3.7 and will be removed two minor releases later. Use ``matplotlib.colormaps[name]`` or ``matplotlib.colormaps.get_cmap(obj)`` instead.

```
plt.scatter(df['Feature 1'], df['Feature 2'], c=y,
cmap=plt.cm.get_cmap('rainbow', len(np.unique(y))))
```



```
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
  warnings.warn(
n_estimators AdaBoost = 1:
Train Accuracy: 0.75, Test Accuracy: 0.72
```

Train Precision: 0.75, Test Precision: 0.72
Train Recall: 0.75, Test Recall: 0.72
Train F1-score: 0.75, Test F1-score: 0.72

n_estimators AdaBoost = 2:
Train Accuracy: 0.83, Test Accuracy: 0.81
Train Precision: 0.83, Test Precision: 0.81
Train Recall: 0.83, Test Recall: 0.81
Train F1-score: 0.83, Test F1-score: 0.81

n_estimators AdaBoost = 5:
Train Accuracy: 0.97, Test Accuracy: 0.97
Train Precision: 0.97, Test Precision: 0.97
Train Recall: 0.97, Test Recall: 0.97
Train F1-score: 0.97, Test F1-score: 0.97

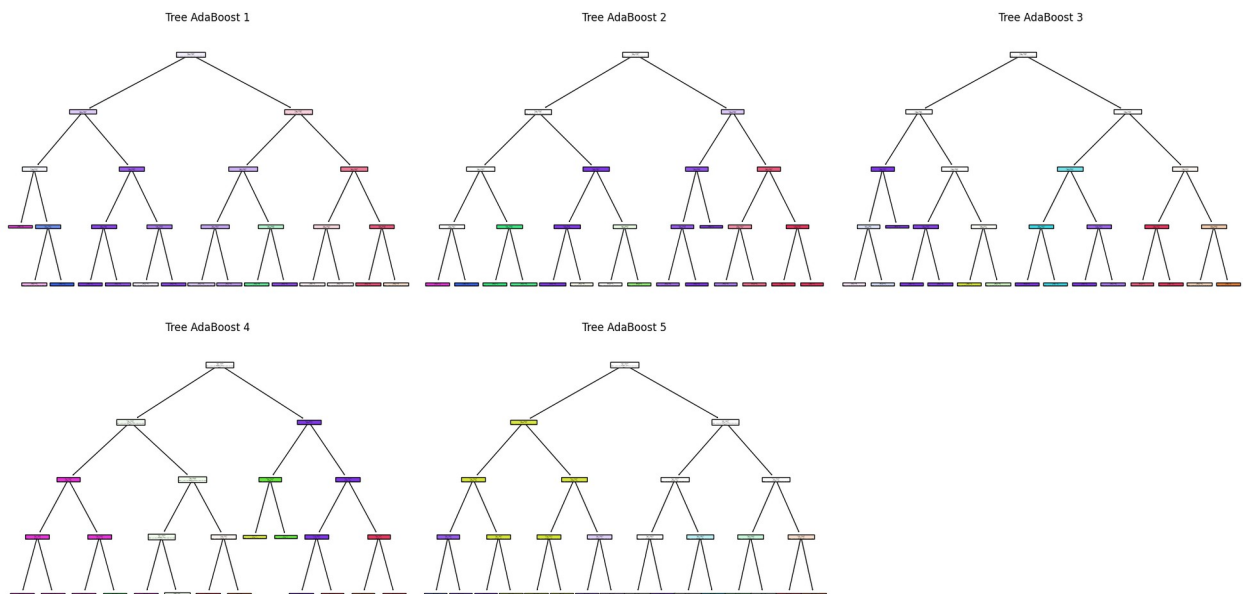

n_estimators AdaBoost = 10:
Train Accuracy: 0.95, Test Accuracy: 0.96
Train Precision: 0.95, Test Precision: 0.96
Train Recall: 0.95, Test Recall: 0.96
Train F1-score: 0.95, Test F1-score: 0.96

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
warnings.warn(

n_estimators AdaBoost = 50:
Train Accuracy: 1.00, Test Accuracy: 1.00
Train Precision: 1.00, Test Precision: 1.00
Train Recall: 1.00, Test Recall: 1.00
Train F1-score: 1.00, Test F1-score: 1.00

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
warnings.warn(

n_estimators AdaBoost = 100:
Train Accuracy: 1.00, Test Accuracy: 0.99
Train Precision: 1.00, Test Precision: 0.99
Train Recall: 1.00, Test Recall: 0.99
Train F1-score: 1.00, Test F1-score: 0.99

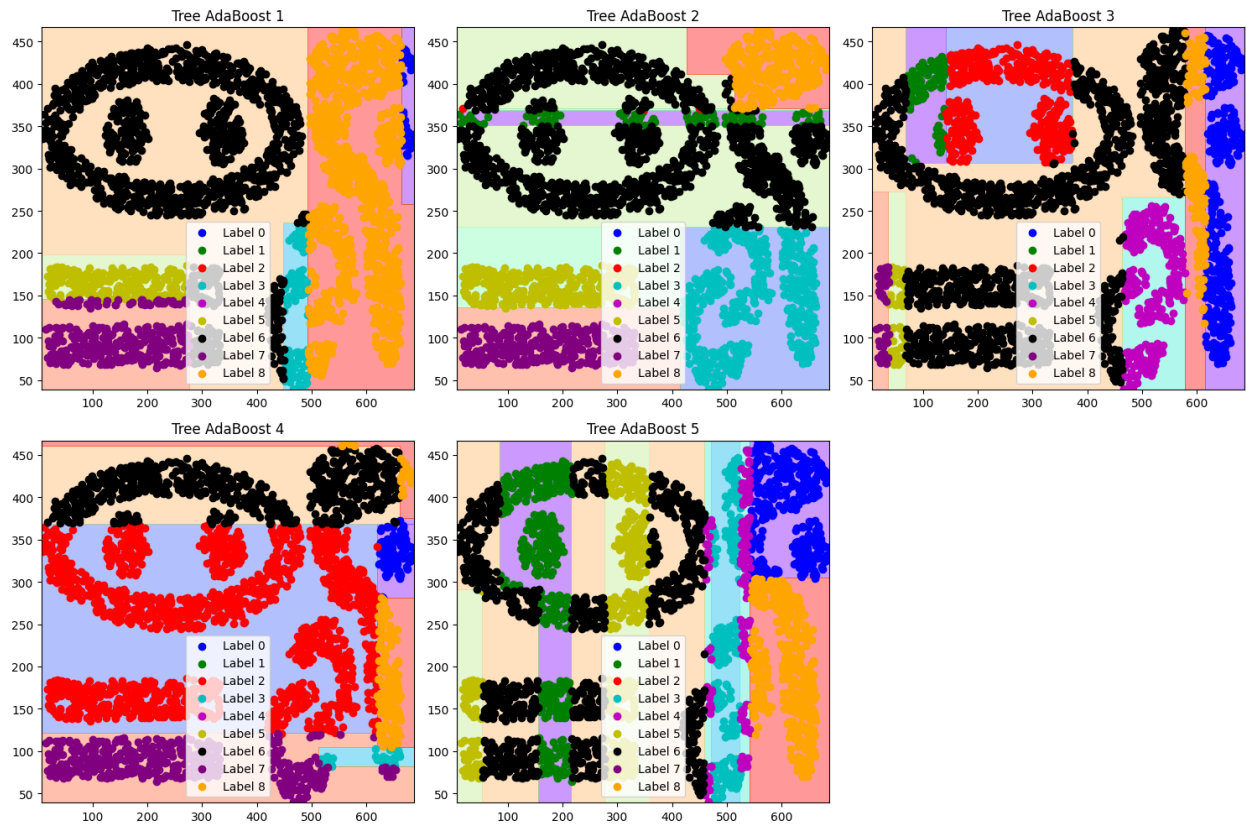


```
<ipython-input-56-9d42626d182c>:116: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432:
UserWarning: X has feature names, but DecisionTreeClassifier was
```

```

fitted without feature names
warnings.warn(
<ipython-input-56-9d42626d182c>:116: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432:
UserWarning: X has feature names, but DecisionTreeClassifier was
fitted without feature names
warnings.warn(
<ipython-input-56-9d42626d182c>:116: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432:
UserWarning: X has feature names, but DecisionTreeClassifier was
fitted without feature names
warnings.warn(
<ipython-input-56-9d42626d182c>:116: MatplotlibDeprecationWarning: The
get_cmap function was deprecated in Matplotlib 3.7 and will be removed
two minor releases later. Use ``matplotlib.colormaps[name]`` or
``matplotlib.colormaps.get_cmap(obj)`` instead.
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.get_cmap('rainbow',
len(np.unique(y))))
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432:
UserWarning: X has feature names, but DecisionTreeClassifier was
fitted without feature names
warnings.warn(

```



Stacked Learning

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier,
AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

# Load the dataset
df = pd.read_csv('/content/drive/My Drive/1/Dataset1.csv')

X = df.drop('Label', axis=1)
y = df['Label']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Define base classifiers
base_classifiers = {
```

```

    'knn': KNeighborsClassifier(n_neighbors=3),
    'decision_tree': DecisionTreeClassifier(max_depth=4,
random_state=42),
    'random_forest': RandomForestClassifier(n_estimators=10,
max_depth=4, random_state=42),
    'adaboost':
AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=4),
n_estimators=10, random_state=42)
}

# Train base classifiers
for name, clf in base_classifiers.items():
    clf.fit(X_train, y_train)
    y_train_pred = clf.predict(X_train)
    y_test_pred = clf.predict(X_test)
    print(f'{name} - Train Accuracy: {accuracy_score(y_train,
y_train_pred):.2f}, Test Accuracy: {accuracy_score(y_test,
y_test_pred):.2f}')

# Create new features for the meta-classifier
meta_features_train = np.column_stack([clf.predict(X_train) for clf in
base_classifiers.values()])
meta_features_test = np.column_stack([clf.predict(X_test) for clf in
base_classifiers.values()])

# Define and train the meta-classifier
meta_classifier = DecisionTreeClassifier(max_depth=4, random_state=42)
meta_classifier.fit(meta_features_train, y_train)

# Evaluate the stacked model
y_train_pred_meta = meta_classifier.predict(meta_features_train)
y_test_pred_meta = meta_classifier.predict(meta_features_test)

train_accuracy = accuracy_score(y_train, y_train_pred_meta)
test_accuracy = accuracy_score(y_test, y_test_pred_meta)

train_precision = precision_score(y_train, y_train_pred_meta,
average='micro')
test_precision = precision_score(y_test, y_test_pred_meta,
average='micro')

train_recall = recall_score(y_train, y_train_pred_meta,
average='micro')
test_recall = recall_score(y_test, y_test_pred_meta, average='micro')

train_f1 = f1_score(y_train, y_train_pred_meta, average='micro')
test_f1 = f1_score(y_test, y_test_pred_meta, average='micro')

print(f'Stacked Model:')
print(f'Train Accuracy: {train_accuracy:.2f}, Test Accuracy:

```



```

{test_accuracy:.2f}')
print(f'Train Precision: {train_precision:.2f}, Test Precision:
{test_precision:.2f}')
print(f'Train Recall: {train_recall:.2f}, Test Recall:
{test_recall:.2f}')
print(f'Train F1-score: {train_f1:.2f}, Test F1-score: {test_f1:.2f}')

knn - Train Accuracy: 1.00, Test Accuracy: 1.00
decision_tree - Train Accuracy: 0.91, Test Accuracy: 0.89
random_forest - Train Accuracy: 0.92, Test Accuracy: 0.90
adaboost - Train Accuracy: 1.00, Test Accuracy: 0.99

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
  warnings.warn(

Stacked Model:
Train Accuracy: 0.90, Test Accuracy: 0.89
Train Precision: 0.90, Test Precision: 0.89
Train Recall: 0.90, Test Recall: 0.89
Train F1-score: 0.90, Test F1-score: 0.89

```

Phase 2

Preprocess

```

import pandas as pd
df = pd.read_csv('/content/drive/My Drive/1/Dataset2.csv')

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4240 entries, 0 to 4239
Data columns (total 16 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   male                  4240 non-null   int64
 1   age                   4240 non-null   int64
 2   education             4135 non-null   float64
 3   currentSmoker         4240 non-null   int64
 4   cigsPerDay            4211 non-null   float64
 5   BPMeds                4187 non-null   float64
 6   prevalentStroke       4240 non-null   int64
 7   prevalentHyp          4240 non-null   int64
 8   diabetes              4240 non-null   int64
 9   totChol               4190 non-null   float64
10  sysBP                4240 non-null   float64

```

```

11  diaBP          4240 non-null  float64
12  BMI           4221 non-null  float64
13  heartRate     4239 non-null  float64
14  glucose       3852 non-null  float64
15  TenYearCHD    4240 non-null  int64
dtypes: float64(9), int64(7)
memory usage: 530.1 KB

df['education'].unique()

array([ 4.,  2.,  1.,  3., nan])

# Calculate the mode of the column
mode_value = df['education'].mode()[0]

# Fill NaN values with the mode
df['education'].fillna(mode_value, inplace=True)

df['cigsPerDay'].unique()

array([ 0., 20., 30., 23., 15.,  9., 10.,  5., 35., 43.,  1., 40.,
  3.,
        2., nan, 12.,  4., 18., 25., 60., 14., 45.,  8., 50., 13.,
 11.,
        7.,  6., 38., 29., 17., 16., 19., 70.])

# Calculate the median of the column
median_value = df['cigsPerDay'].median()

# Fill NaN values with the median
df['cigsPerDay'].fillna(median_value, inplace=True)

df['BPMeds'].unique()

array([ 0.,  1., nan])

df = df.dropna(subset=['BPMeds'])

# Calculate the median of the column
median_value = df['totChol'].median()

# Fill NaN values with the median
df['totChol'].fillna(median_value, inplace=True)

<ipython-input-66-6f836d55e12c>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#
returning-a-view-versus-a-copy
df['totChol'].fillna(median_value, inplace=True)

```

```

# Calculate the median of the column
mean_value = df['BMI'].mean()

# Fill NaN values with the median
df['BMI'].fillna(mean_value, inplace=True)

<ipython-input-67-a33f1e4e780a>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
    df['BMI'].fillna(mean_value, inplace=True)

# Calculate the median of the column
median_value = df['heartRate'].median()

# Fill NaN values with the median
df['heartRate'].fillna(median_value, inplace=True)

<ipython-input-68-25274b80e945>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
    df['heartRate'].fillna(median_value, inplace=True)

# Calculate the median of the column
median_value = df['glucose'].median()

# Fill NaN values with the median
df['glucose'].fillna(median_value, inplace=True)

<ipython-input-69-3502970513b1>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
    df['glucose'].fillna(median_value, inplace=True)

df.head()

{"summary": "{\n  \"name\": \"df\",\n  \"rows\": 4187,\n  \"fields\": [\n    {\n      \"column\": \"male\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 1,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          0,\n          1\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"age\",\n      \"properties\": {\n        \"dtype\": \"

```

```

\"number\", \n          \"std\": 8, \n          \"min\": 32, \n
\"max\": 70, \n          \"num_unique_values\": 39, \n          \"samples\":
[\n          34, \n          70 \n          ], \n          \"semantic_type\":
\"\", \n          \"description\": \"\" \n          } \n          }, \n          { \n
\"column\": \"education\", \n          \"properties\": { \n
\"dtype\": \"number\", \n          \"std\": 1.01787672260344, \n
\"min\": 1.0, \n          \"max\": 4.0, \n          \"num_unique_values\":
4, \n          \"samples\": [\n          2.0, \n          3.0 \n          ], \n
          \"semantic_type\": \"\", \n          \"description\": \"\" \n
          } \n          }, \n          { \n          \"column\": \"currentSmoker\", \n
\"properties\": { \n          \"dtype\": \"number\", \n          \"std\":
0, \n          \"min\": 0, \n          \"max\": 1, \n
\"num_unique_values\": 2, \n          \"samples\": [\n          1, \n
0 \n          ], \n          \"semantic_type\": \"\", \n
          \"description\": \"\" \n          } \n          }, \n          { \n          \"column\":
\"cigsPerDay\", \n          \"properties\": { \n          \"dtype\":
\"number\", \n          \"std\": 11.901904311142495, \n          \"min\":
0.0, \n          \"max\": 70.0, \n          \"num_unique_values\": 33, \n
          \"samples\": [\n          19.0, \n          4.0 \n          ], \n
          \"semantic_type\": \"\", \n          \"description\": \"\" \n
          } \n          }, \n          { \n          \"column\": \"BPMeds\", \n
          \"properties\": { \n          \"dtype\": \"number\", \n          \"std\":
0.16954428739626196, \n          \"min\": 0.0, \n          \"max\": 1.0, \n
          \"num_unique_values\": 2, \n          \"samples\": [\n          1.0, \n
0.0 \n          ], \n          \"semantic_type\": \"\", \n
          \"description\": \"\" \n          } \n          }, \n          { \n          \"column\":
\"prevalentStroke\", \n          \"properties\": { \n          \"dtype\":
\"number\", \n          \"std\": 0, \n          \"min\": 0, \n
          \"max\": 1, \n          \"num_unique_values\": 2, \n          \"samples\":
[\n          1, \n          0 \n          ], \n          \"semantic_type\":
\"\", \n          \"description\": \"\" \n          } \n          }, \n          { \n
          \"column\": \"prevalentHyp\", \n          \"properties\": { \n
          \"dtype\": \"number\", \n          \"std\": 0, \n          \"min\": 0, \n
          \"max\": 1, \n          \"num_unique_values\": 2, \n          \"samples\":
[\n          1, \n          0 \n          ], \n          \"semantic_type\":
\"\", \n          \"description\": \"\" \n          } \n          }, \n          { \n
          \"column\": \"diabetes\", \n          \"properties\": { \n
          \"dtype\": \"number\", \n          \"std\": 0, \n          \"min\": 0, \n
          \"max\": 1, \n          \"num_unique_values\": 2, \n          \"samples\":
[\n          1, \n          0 \n          ], \n          \"semantic_type\":
\"\", \n          \"description\": \"\" \n          } \n          }, \n          { \n
          \"column\": \"totChol\", \n          \"properties\": { \n
          \"dtype\": \"number\", \n          \"std\": 44.2237265788299, \n          \"min\":
107.0, \n          \"max\": 696.0, \n          \"num_unique_values\": 248, \n
          \"samples\": [\n          311.0, \n          205.0 \n          ], \n
          \"semantic_type\": \"\", \n          \"description\": \"\" \n
          } \n          }, \n          { \n          \"column\": \"sysBP\", \n
          \"properties\": { \n          \"dtype\": \"number\", \n          \"std\":
21.98290660966593, \n          \"min\": 83.5, \n          \"max\": 295.0, \n

```

```

{"num_unique_values": 233, "samples": [117.0, 191.5], "semantic_type": "\"", "description": "\"", "column": "diaBP", "properties": {"dtype": "number", "std": 11.878463990074186, "min": 48.0, "max": 142.5, "num_unique_values": 146, "samples": [86.5, 108.5], "semantic_type": "\"", "description": "\"", "column": "BMI", "properties": {"dtype": "number", "std": 4.067178635883306, "min": 15.54, "max": 56.8, "num_unique_values": 1360, "samples": [19.66, 30.99], "semantic_type": "\"", "description": "\"", "column": "heartRate", "properties": {"dtype": "number", "std": 12.053463897069951, "min": 44.0, "max": 143.0, "num_unique_values": 73, "samples": [47.0, 85.0], "semantic_type": "\"", "description": "\"", "column": "glucose", "properties": {"dtype": "number", "std": 22.932432152397404, "min": 40.0, "max": 394.0, "num_unique_values": 142, "samples": [332.0, 74.0], "semantic_type": "\"", "description": "\"", "column": "TenYearCHD", "properties": {"dtype": "number", "std": 0, "min": 0, "max": 1, "num_unique_values": 2, "samples": [0, 1], "semantic_type": "\"", "description": "\""}], "type": "dataframe", "variable name": "df"}

```

```
# Calculate the correlation matrix
```

```
correlation matrix = df.corr()
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(10, 8))
```


```
# Draw the heatmap
```

```
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', vmin=-1,
vmax=1, center=0)
```

```
# Add title and labels
```

```
plt.title('Correlation Matrix Heatmap')
```

```
plt.show()
```



```
columns_to_drop = set()

# Iterate through the correlation matrix
for i in range(len(correlation_matrix.columns)):
    for j in range(i):
        if abs(correlation_matrix.iloc[i, j]) > threshold:
            column_to_drop = correlation_matrix.columns[i]
            columns_to_drop.add(column_to_drop)

# Create a new DataFrame with only the columns that have low correlation
df_low_corr = df.drop(columns=columns_to_drop)

print("Columns to drop due to high correlation:")
print(columns_to_drop)
print("\nDataFrame with only low correlated columns:")
print(df_low_corr)
```

Columns to drop due to high correlation:
{'cigsPerDay', 'diaBP'}

DataFrame with only low correlated columns:

	male	age	education	currentSmoker	BPMeds	prevalentStroke	\
0	1	39	4.0	0	0.0	0	
1	0	46	2.0	0	0.0	0	
2	1	48	1.0	1	0.0	0	
3	0	61	3.0	1	0.0	0	
4	0	46	3.0	1	0.0	0	
...
4234	1	51	3.0	1	0.0	0	
4236	0	44	1.0	1	0.0	0	
4237	0	52	2.0	0	0.0	0	
4238	1	40	3.0	0	0.0	0	
4239	0	39	3.0	1	0.0	0	

	prevalentHyp	diabetes	totChol	sysBP	BMI	heartRate	glucose	\
0	0	0	195.0	106.0	26.97	80.0	77.0	
1	0	0	250.0	121.0	28.73	95.0	76.0	
2	0	0	245.0	127.5	25.34	75.0	70.0	
3	1	0	225.0	150.0	28.58	65.0	103.0	
4	0	0	285.0	130.0	23.10	85.0	85.0	
...
4234	0	0	207.0	126.5	19.71	65.0	68.0	
4236	0	0	210.0	126.5	19.16	86.0	78.0	
4237	0	0	269.0	133.5	21.47	80.0	107.0	
4238	1	0	185.0	141.0	25.60	67.0	72.0	
4239	0	0	196.0	133.0	20.91	85.0	80.0	

	TenYearCHD
0	0
1	0
2	0
3	1
4	0
...	...
4234	0

4236	0
4237	0
4238	0
4239	0

[4187 rows x 14 columns]

df_low_corr

```
{"summary":{"\n  \"name\": \"df_low_corr\",\n  \"rows\": 4187,\n  \"fields\": [\n    {\n      \"column\": \"male\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 1,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          0,\n          1\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"age\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 8,\n        \"min\": 32,\n        \"max\": 70,\n        \"num_unique_values\": 39,\n        \"samples\": [\n          34,\n          70\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"education\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 1.01787672260344,\n        \"min\": 1.0,\n        \"max\": 4.0,\n        \"num_unique_values\": 4,\n        \"samples\": [\n          2.0,\n          3.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"currentSmoker\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 1,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          0,\n          1\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"BPMeds\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.16954428739626196,\n        \"min\": 0.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          0.0,\n          1.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"prevalentStroke\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 1,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          0,\n          1\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"prevalentHyp\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 1,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          0,\n          1\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"diabetes\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0,\n        \"min\": 0,\n        \"max\": 1,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          0,\n          1\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    ]\n  }\n}
```



```
[\n          1,\n          0\n          ],\n          \"semantic_type\":  
\"\",\n          \"description\": \"\"\n          },\n          {\n          \"column\": \"totChol\", \n          \"properties\": {\n          \"dtype\":  
\"number\", \n          \"std\": 44.2237265788299, \n          \"min\":  
107.0, \n          \"max\": 696.0, \n          \"num_unique_values\": 248, \n          \"samples\": [\n          311.0, \n          205.0\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }, \n          {\n          \"column\": \"sysBP\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 21.98290660966593, \n          \"min\": 83.5, \n          \"max\": 295.0, \n          \"num_unique_values\": 233, \n          \"samples\": [\n          117.0, \n          191.5\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }, \n          {\n          \"column\":  
\"BMI\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 4.067178635883306, \n          \"min\": 15.54, \n          \"max\":  
56.8, \n          \"num_unique_values\": 1360, \n          \"samples\": [\n          19.66, \n          30.99\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }, \n          {\n          \"column\": \"heartRate\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 12.053463897069951, \n          \"min\": 44.0, \n          \"max\": 143.0, \n          \"num_unique_values\": 73, \n          \"samples\": [\n          85.0, \n          47.0\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }, \n          {\n          \"column\":  
\"glucose\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 22.932432152397404, \n          \"min\": 40.0, \n          \"max\": 394.0, \n          \"num_unique_values\": 142, \n          \"samples\": [\n          332.0, \n          74.0\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }, \n          {\n          \"column\": \"TenYearCHD\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\":  
0, \n          \"min\": 0, \n          \"max\": 1, \n          \"num_unique_values\": 2, \n          \"samples\": [\n          1, \n          0\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n          }\n          }\n          ], \n          \"type\": \"dataframe\", \"variable_name\": \"df_low_corr\"}
```

Bagging

```
# Define 9 different colors
colors = ['b', 'g', 'r', 'c', 'm', 'y', 'k', 'purple', 'orange']

X = df_low_corr.drop('TenYearCHD', axis=1)
y = df_low_corr['TenYearCHD']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```

# Number of estimators to evaluate
num_estimators = [1, 2, 5, 10, 50, 100]

train_accuracies = []
test_accuracies = []

# Loop to test different numbers of estimators
for n in num_estimators:
    # Create Bagging model with different numbers of estimators
    bagging_clf =
BaggingClassifier(base_estimator=DecisionTreeClassifier(max_depth=4,
splitter='random'),
                    n_estimators=n,
                    random_state=42)

    # Train the model on training data
    bagging_clf.fit(X_train, y_train)

    # Predict labels for training and testing data
    y_train_pred = bagging_clf.predict(X_train)
    y_test_pred = bagging_clf.predict(X_test)

    # Compute model accuracy
    train_accuracy = accuracy_score(y_train, y_train_pred)
    test_accuracy = accuracy_score(y_test, y_test_pred)

    # Store accuracies in lists
    train_accuracies.append(train_accuracy)
    test_accuracies.append(test_accuracy)

    # Evaluate the model for training and testing data
    train_f1 = f1_score(y_train, y_train_pred, average='micro')
    test_f1 = f1_score(y_test, y_test_pred, average='micro')

    train_precision = precision_score(y_train, y_train_pred,
average='micro')
    test_precision = precision_score(y_test, y_test_pred,
average='micro')

    train_recall = recall_score(y_train, y_train_pred,
average='micro')
    test_recall = recall_score(y_test, y_test_pred, average='micro')

    print(f'n_estimators Bagging = {n}:')
    print(f'Train Accuracy: {train_accuracy:.2f}, Test Accuracy:
{test_accuracy:.2f}')
    print(f'Train Precision: {train_precision:.2f}, Test Precision:
{test_precision:.2f}')
    print(f'Train Recall: {train_recall:.2f}, Test Recall:
{test_recall:.2f}')

```

```

    print(f'Train F1-score: {train_f1:.2f}, Test F1-score:
{test_f1:.2f}')
    print()

# Plot Bagging accuracy versus number of estimators
plt.figure(figsize=(10, 6))
plt.plot(num_estimators, train_accuracies, marker='o', label='Train
Accuracy')
plt.plot(num_estimators, test_accuracies, marker='o', label='Test
Accuracy')
plt.title('Bagging Accuracy vs. Number of Estimators')
plt.xlabel('Number of Estimators')
plt.ylabel('Accuracy')
plt.xticks(num_estimators)
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

```

```

n_estimators Bagging = 1:
Train Accuracy: 0.86, Test Accuracy: 0.83
Train Precision: 0.86, Test Precision: 0.83
Train Recall: 0.86, Test Recall: 0.83
Train F1-score: 0.86, Test F1-score: 0.83

```

```

n_estimators Bagging = 2:
Train Accuracy: 0.86, Test Accuracy: 0.83
Train Precision: 0.86, Test Precision: 0.83
Train Recall: 0.86, Test Recall: 0.83
Train F1-score: 0.86, Test F1-score: 0.83

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version

```

1.2 and will be removed in 1.4.

```
warnings.warn(
```

```
n_estimators Bagging = 5:
```

```
Train Accuracy: 0.86, Test Accuracy: 0.83
```

```
Train Precision: 0.86, Test Precision: 0.83
```

```
Train Recall: 0.86, Test Recall: 0.83
```

```
Train F1-score: 0.86, Test F1-score: 0.83
```

```
n_estimators Bagging = 10:
```

```
Train Accuracy: 0.86, Test Accuracy: 0.83
```

```
Train Precision: 0.86, Test Precision: 0.83
```

```
Train Recall: 0.86, Test Recall: 0.83
```

```
Train F1-score: 0.86, Test F1-score: 0.83
```

```
n_estimators Bagging = 50:
```

```
Train Accuracy: 0.86, Test Accuracy: 0.83
```

```
Train Precision: 0.86, Test Precision: 0.83
```

```
Train Recall: 0.86, Test Recall: 0.83
```

```
Train F1-score: 0.86, Test F1-score: 0.83
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
```

```
FutureWarning: `base_estimator` was renamed to `estimator` in version  
1.2 and will be removed in 1.4.
```

```
warnings.warn(
```

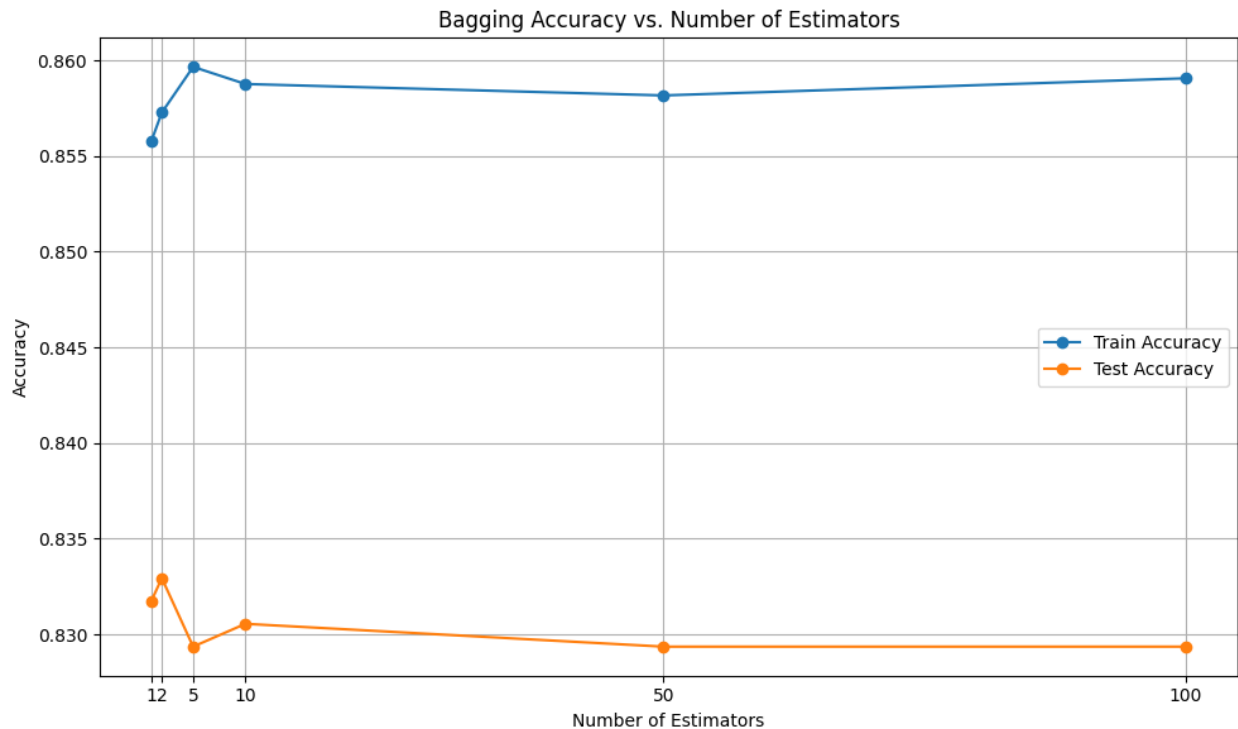
```
n_estimators Bagging = 100:
```

```
Train Accuracy: 0.86, Test Accuracy: 0.83
```

```
Train Precision: 0.86, Test Precision: 0.83
```

```
Train Recall: 0.86, Test Recall: 0.83
```

```
Train F1-score: 0.86, Test F1-score: 0.83
```



RandomForest

```
# Define 9 different colors
colors = ['b', 'g', 'r', 'c', 'm', 'y', 'k', 'purple', 'orange']

X = df_low_corr.drop('TenYearCHD', axis=1)
y = df_low_corr['TenYearCHD']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Number of estimators to evaluate
num_estimators = [1, 2, 5, 10, 50, 100]

train_accuracies = []
test_accuracies = []

# Loop to test different numbers of estimators
for n in num_estimators:
    # Create RandomForest model with different numbers of estimators
    rf_clf = RandomForestClassifier(n_estimators=n, max_depth=4,
random_state=42)
```

```

# Train the model on training data
rf_clf.fit(X_train, y_train)

# Predict labels for training and testing data
y_train_pred = rf_clf.predict(X_train)
y_test_pred = rf_clf.predict(X_test)

# Compute model accuracy
train_accuracy = accuracy_score(y_train, y_train_pred)
test_accuracy = accuracy_score(y_test, y_test_pred)

# Store accuracies in lists
train_accuracies.append(train_accuracy)
test_accuracies.append(test_accuracy)

# Evaluate the model for training and testing data
train_f1 = f1_score(y_train, y_train_pred, average='micro')
test_f1 = f1_score(y_test, y_test_pred, average='micro')

train_precision = precision_score(y_train, y_train_pred,
average='micro')
test_precision = precision_score(y_test, y_test_pred,
average='micro')

train_recall = recall_score(y_train, y_train_pred,
average='micro')
test_recall = recall_score(y_test, y_test_pred, average='micro')

print(f'n_estimators RandomForest = {n}:')
print(f'Train Accuracy: {train_accuracy:.2f}, Test Accuracy:
{test_accuracy:.2f}')
print(f'Train Precision: {train_precision:.2f}, Test Precision:
{test_precision:.2f}')
print(f'Train Recall: {train_recall:.2f}, Test Recall:
{test_recall:.2f}')
print(f'Train F1-score: {train_f1:.2f}, Test F1-score:
{test_f1:.2f}')
print()

# Plot RandomForest accuracy versus number of estimators
plt.figure(figsize=(10, 6))
plt.plot(num_estimators, train_accuracies, marker='o', label='Train
Accuracy')
plt.plot(num_estimators, test_accuracies, marker='o', label='Test
Accuracy')
plt.title('RandomForest Accuracy vs. Number of Estimators')
plt.xlabel('Number of Estimators')
plt.ylabel('Accuracy')
plt.xticks(num_estimators)
plt.legend()

```

```
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
n_estimators RandomForest = 1:
Train Accuracy: 0.86, Test Accuracy: 0.83
Train Precision: 0.86, Test Precision: 0.83
Train Recall: 0.86, Test Recall: 0.83
Train F1-score: 0.86, Test F1-score: 0.83
```

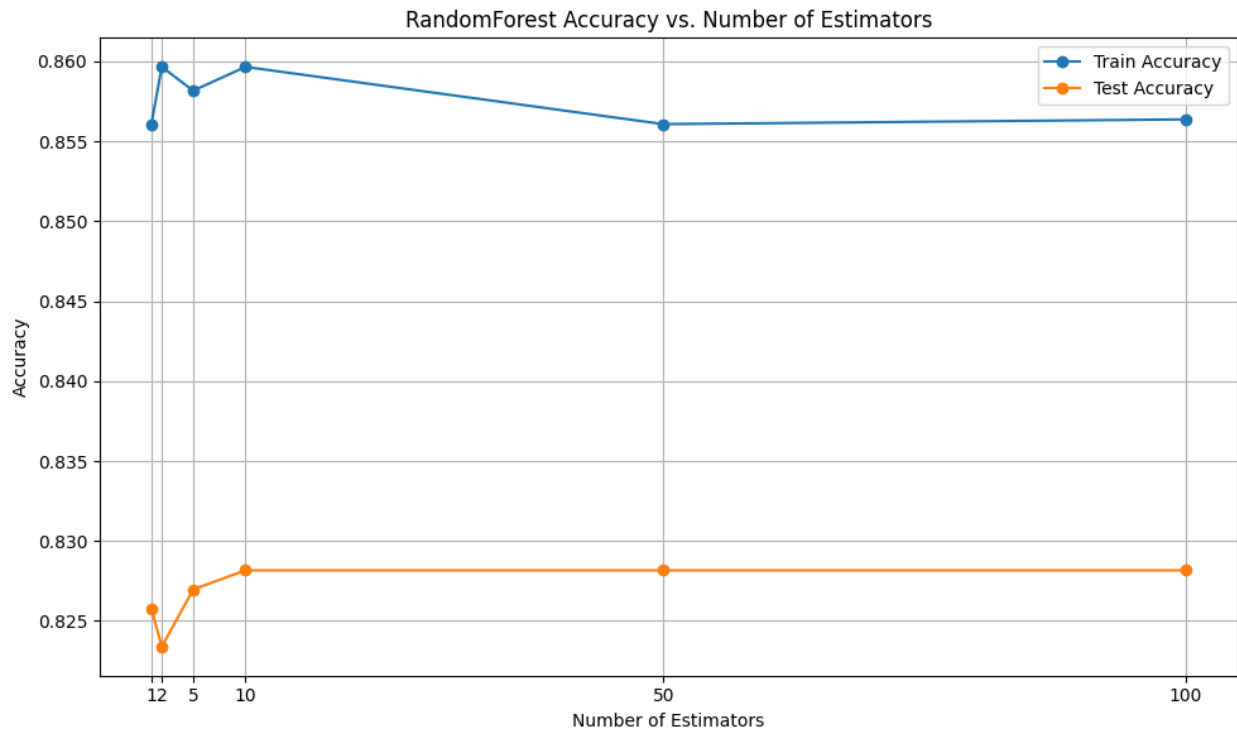
```
n_estimators RandomForest = 2:
Train Accuracy: 0.86, Test Accuracy: 0.82
Train Precision: 0.86, Test Precision: 0.82
Train Recall: 0.86, Test Recall: 0.82
Train F1-score: 0.86, Test F1-score: 0.82
```

```
n_estimators RandomForest = 5:
Train Accuracy: 0.86, Test Accuracy: 0.83
Train Precision: 0.86, Test Precision: 0.83
Train Recall: 0.86, Test Recall: 0.83
Train F1-score: 0.86, Test F1-score: 0.83
```

```
n_estimators RandomForest = 10:
Train Accuracy: 0.86, Test Accuracy: 0.83
Train Precision: 0.86, Test Precision: 0.83
Train Recall: 0.86, Test Recall: 0.83
Train F1-score: 0.86, Test F1-score: 0.83
```

```
n_estimators RandomForest = 50:
Train Accuracy: 0.86, Test Accuracy: 0.83
Train Precision: 0.86, Test Precision: 0.83
Train Recall: 0.86, Test Recall: 0.83
Train F1-score: 0.86, Test F1-score: 0.83
```

```
n_estimators RandomForest = 100:
Train Accuracy: 0.86, Test Accuracy: 0.83
Train Precision: 0.86, Test Precision: 0.83
Train Recall: 0.86, Test Recall: 0.83
Train F1-score: 0.86, Test F1-score: 0.83
```



Adaboost

```
# Define 9 different colors
colors = ['b', 'g', 'r', 'c', 'm', 'y', 'k', 'purple', 'orange']

X = df_low_corr.drop('TenYearCHD', axis=1)
y = df_low_corr['TenYearCHD']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Number of estimators to evaluate
num_estimators = [1, 2, 5, 10, 50, 100]

train_accuracies = []
test_accuracies = []

# Loop to test different numbers of estimators
for n in num_estimators:
    # Create AdaBoost model with different numbers of estimators
    ada_clf =
AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=4,
splitter='random'),
```



```

n_estimators=n, random_state=42)

# Train the model on training data
ada_clf.fit(X_train, y_train)

# Predict labels for training and testing data
y_train_pred = ada_clf.predict(X_train)
y_test_pred = ada_clf.predict(X_test)

# Compute model accuracy
train_accuracy = accuracy_score(y_train, y_train_pred)
test_accuracy = accuracy_score(y_test, y_test_pred)

# Store accuracies in lists
train_accuracies.append(train_accuracy)
test_accuracies.append(test_accuracy)

# Evaluate the model for training and testing data
train_f1 = f1_score(y_train, y_train_pred, average='micro')
test_f1 = f1_score(y_test, y_test_pred, average='micro')

train_precision = precision_score(y_train, y_train_pred,
average='micro')
test_precision = precision_score(y_test, y_test_pred,
average='micro')

train_recall = recall_score(y_train, y_train_pred,
average='micro')
test_recall = recall_score(y_test, y_test_pred, average='micro')

print(f'n_estimators AdaBoost = {n}:')
print(f'Train Accuracy: {train_accuracy:.2f}, Test Accuracy:
{test_accuracy:.2f}')
print(f'Train Precision: {train_precision:.2f}, Test Precision:
{test_precision:.2f}')
print(f'Train Recall: {train_recall:.2f}, Test Recall:
{test_recall:.2f}')
print(f'Train F1-score: {train_f1:.2f}, Test F1-score:
{test_f1:.2f}')
print()

# Plot AdaBoost accuracy versus number of estimators
plt.figure(figsize=(10, 6))
plt.plot(num_estimators, train_accuracies, marker='o', label='Train
Accuracy')
plt.plot(num_estimators, test_accuracies, marker='o', label='Test
Accuracy')
plt.title('AdaBoost Accuracy vs. Number of Estimators')
plt.xlabel('Number of Estimators')
plt.ylabel('Accuracy')

```

```
plt.xticks(num_estimators)
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
```

```
warnings.warn(
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
```

```
warnings.warn(
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
```

```
warnings.warn(
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
```

```
warnings.warn(
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
```

```
warnings.warn(
```

```
n_estimators AdaBoost = 1:
Train Accuracy: 0.86, Test Accuracy: 0.82
Train Precision: 0.86, Test Precision: 0.82
Train Recall: 0.86, Test Recall: 0.82
Train F1-score: 0.86, Test F1-score: 0.82
```

```
n_estimators AdaBoost = 2:
Train Accuracy: 0.86, Test Accuracy: 0.82
Train Precision: 0.86, Test Precision: 0.82
Train Recall: 0.86, Test Recall: 0.82
Train F1-score: 0.86, Test F1-score: 0.82
```

```
n_estimators AdaBoost = 5:
Train Accuracy: 0.86, Test Accuracy: 0.82
Train Precision: 0.86, Test Precision: 0.82
Train Recall: 0.86, Test Recall: 0.82
Train F1-score: 0.86, Test F1-score: 0.82
```

```
n_estimators AdaBoost = 10:
Train Accuracy: 0.87, Test Accuracy: 0.81
Train Precision: 0.87, Test Precision: 0.81
Train Recall: 0.87, Test Recall: 0.81
```

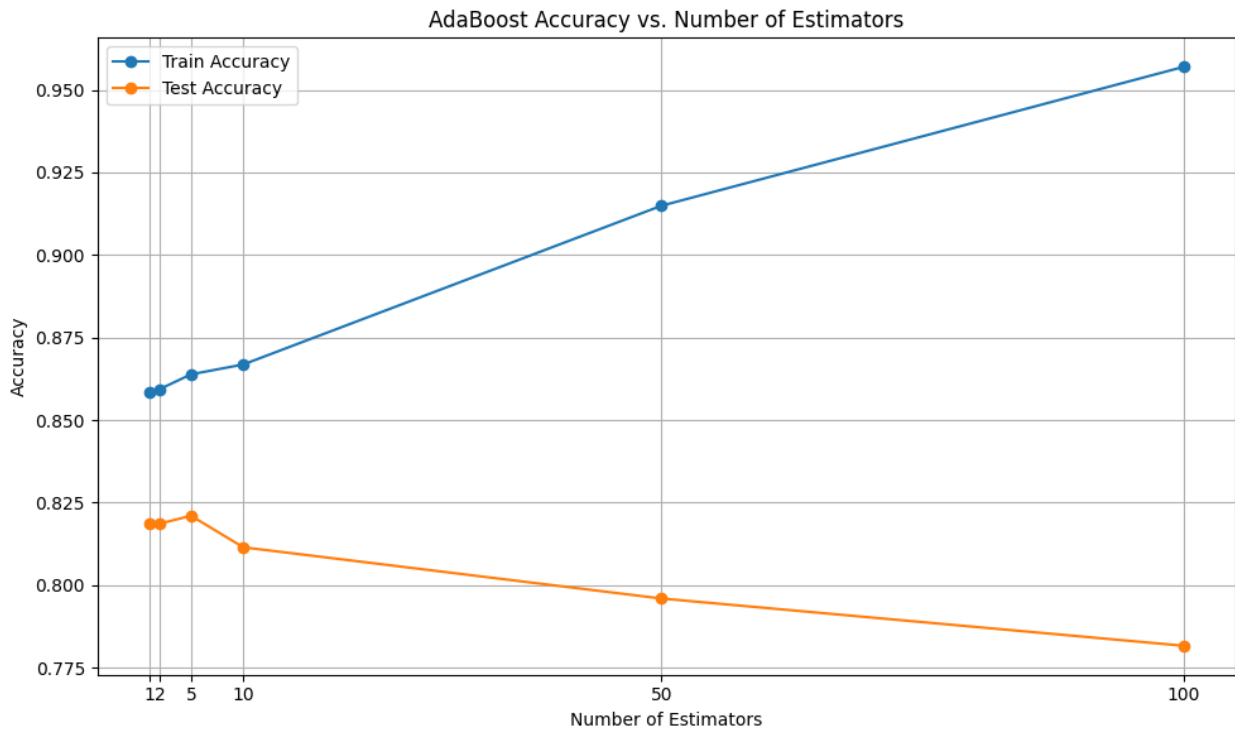
Train F1-score: 0.87, Test F1-score: 0.81

```
n_estimators AdaBoost = 50:  
Train Accuracy: 0.91, Test Accuracy: 0.80  
Train Precision: 0.91, Test Precision: 0.80  
Train Recall: 0.91, Test Recall: 0.80  
Train F1-score: 0.91, Test F1-score: 0.80
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:  
FutureWarning: `base_estimator` was renamed to `estimator` in version  
1.2 and will be removed in 1.4.  
warnings.warn(  

```

```
n_estimators AdaBoost = 100:  
Train Accuracy: 0.96, Test Accuracy: 0.78  
Train Precision: 0.96, Test Precision: 0.78  
Train Recall: 0.96, Test Recall: 0.78  
Train F1-score: 0.96, Test F1-score: 0.78
```



Stacked Learner

```
import numpy as np  
import pandas as pd
```

```

import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier,
AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

X = df_low_corr.drop('TenYearCHD', axis=1)
y = df_low_corr['TenYearCHD']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Define base classifiers
base_classifiers = {
    'knn': KNeighborsClassifier(n_neighbors=3),
    'decision_tree': DecisionTreeClassifier(max_depth=4,
random_state=42),
    'random_forest': RandomForestClassifier(n_estimators=10,
max_depth=4, random_state=42),
    'adaboost':
AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=4),
n_estimators=10, random_state=42)
}

# Train base classifiers
for name, clf in base_classifiers.items():
    clf.fit(X_train, y_train)
    y_train_pred = clf.predict(X_train)
    y_test_pred = clf.predict(X_test)
    print(f'{name} - Train Accuracy: {accuracy_score(y_train,
y_train_pred):.2f}, Test Accuracy: {accuracy_score(y_test,
y_test_pred):.2f}')

# Create new features for the meta-classifier
meta_features_train = np.column_stack([clf.predict(X_train) for clf in
base_classifiers.values()])
meta_features_test = np.column_stack([clf.predict(X_test) for clf in
base_classifiers.values()])

# Define and train the meta-classifier
meta_classifier = DecisionTreeClassifier(max_depth=4, random_state=42)
meta_classifier.fit(meta_features_train, y_train)

# Evaluate the stacked model
y_train_pred_meta = meta_classifier.predict(meta_features_train)
y_test_pred_meta = meta_classifier.predict(meta_features_test)

```

```

train_accuracy = accuracy_score(y_train, y_train_pred_meta)
test_accuracy = accuracy_score(y_test, y_test_pred_meta)

train_precision = precision_score(y_train, y_train_pred_meta,
average='micro')
test_precision = precision_score(y_test, y_test_pred_meta,
average='micro')

train_recall = recall_score(y_train, y_train_pred_meta,
average='micro')
test_recall = recall_score(y_test, y_test_pred_meta, average='micro')

train_f1 = f1_score(y_train, y_train_pred_meta, average='micro')
test_f1 = f1_score(y_test, y_test_pred_meta, average='micro')

print(f'Stacked Model:')
print(f'Train Accuracy: {train_accuracy:.2f}, Test Accuracy:
{test_accuracy:.2f}')
print(f'Train Precision: {train_precision:.2f}, Test Precision:
{test_precision:.2f}')
print(f'Train Recall: {train_recall:.2f}, Test Recall:
{test_recall:.2f}')
print(f'Train F1-score: {train_f1:.2f}, Test F1-score: {test_f1:.2f}')

knn - Train Accuracy: 0.89, Test Accuracy: 0.79
decision_tree - Train Accuracy: 0.86, Test Accuracy: 0.82
random_forest - Train Accuracy: 0.86, Test Accuracy: 0.83
adaboost - Train Accuracy: 0.88, Test Accuracy: 0.80

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166:
FutureWarning: `base_estimator` was renamed to `estimator` in version
1.2 and will be removed in 1.4.
  warnings.warn(

Stacked Model:
Train Accuracy: 0.89, Test Accuracy: 0.78
Train Precision: 0.89, Test Precision: 0.78
Train Recall: 0.89, Test Recall: 0.78
Train F1-score: 0.89, Test F1-score: 0.78

```